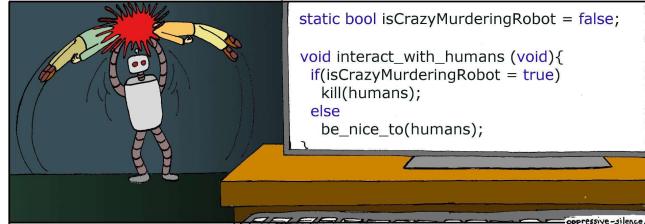


Lecture #16 - That's all folks!

- Intro to Graphs
- Graph Traversals
 - Depth-first
 - Breadth-first
- Dijkstra's Algorithm

Final Exam: Saturday, March 16th 11:30am-2:30pm
Final Exam Location: TBA

Graphs



Graphs Why should you care?

Facebook? Duh!

Not good enough?
Google+?

Computer Animation?

Google Maps?

The Internet?

So pay attention!



Introduction to Graphs

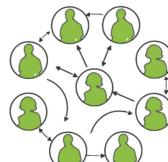
A **graph** is an ADT that stores a set of **entities** and also keeps track of the **relationships** between all of them.

Examples of Entities

People
Cities
Web pages

Examples of Relationships

Joe is friends with Linda
LA is 3000 miles from NYC
ucla.edu links to awesome.com



Introduction to Graphs

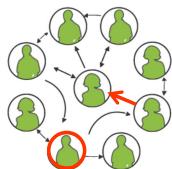
Each graph holds two types of items:

Vertices (aka Nodes):

A vertex might represent a **person**, a **city** or a **web page**.

Edges (aka Arcs):

An edge simply **connects two* vertices** to each other.



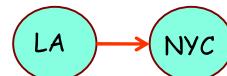
* Technically, an edge could connect a vertex to itself!

Directed vs. Undirected Graphs

There are two major types of graphs...

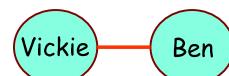
Directed Graph

In a **directed graph**, an edge goes from one vertex to another in a **specific direction**.



Undirected Graph

In an **undirected graph**, all edges are **bi-directional**. You can go either way along any edge.



For example, above we have an edge that goes from the LA vertex to NYC vertex, but not the other way around.

(e.g., there may be a flight from LA to NYC but not the other way around)

For example, Vickie and Ben are mutual friends on FaceBook.

(It would be kinda creepy if Vickie liked Ben, but not visa-versa)

Representing a Graph in Your Programs

The easiest way to represent a graph is with a **double-dimensional array**.

The **size** of both dimensions of the array is equal to the **number of vertices** in the graph.

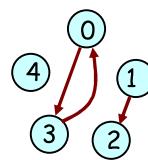
```
bool graph[5][5];
```

Each element in the array indicates whether or not there is an edge between vertex *i* and vertex *j*.

Representing a Graph in Your Programs

Each element in the array indicates whether or not there is an edge between vertex *i* and vertex *j*.

```
bool graph[5][5];
// edge from vertex 0 to vertex 3
graph[0][3] = true;
graph[1][2] = true;
graph[3][0] = true;
```



	0	1	2	3	4
0				T	
1			T		
2					
3	T				
4					

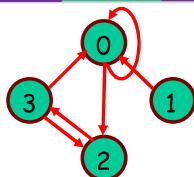
As you can see, when we set array[i][j] to **true**, it represents a **directed edge** from vertex *i* to vertex *j*.

This is called an **adjacency matrix**.

Representing a Graph in Your Programs

Exercise: What does the following directed graph look like?

Nodes	0	1	2	3
0	True	False	True	False
1	True	False	False	False
2	False	False	False	True
3	True	False	True	False



Representing a Graph in Your Programs

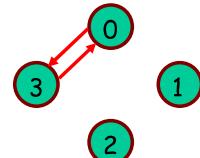
Question:
How do you represent an **undirected graph** with an adjacency matrix?

It's easy!

To bi-directionally connect vertices *i* and *j*, simply set array[i][j] to **true** and set array[j][i] to **true** as well!

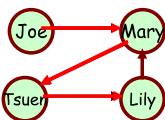
```
graph[0][3] = true;
graph[3][0] = true;
```

Nodes	0	1	2	3
0				True
1				
2				
3	True			



An Interesting Property of Adjacency Matrices

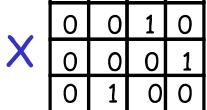
Consider the following graph: And its associated A.M.:



	Joe	Mary	Tsuen	Lily
Joe	0	1	0	0
Mary	0	0	1	0
Tsuen	0	0	0	1
Lily	0	1	0	0

Neato effect: If you multiply the matrix by itself something cool happens!

0	1	0	0
0	0	1	0
0	0	0	1
0	1	0	0



	Joe	Mary	Tsuen	Lily
Joe	0	0	1	0
Mary	0	0	0	1
Tsuen	0	1	0	0
Lily	0	0	1	0

=

0	1	0	0
0	0	1	0
0	0	0	1
0	1	0	0

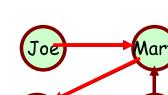
0	1	0	0
0	0	1	0
0	0	0	1
0	1	0	0

	Joe	Mary	Tsuen	Lily
Joe	0	0	0	1
Mary	0	0	0	0
Tsuen	0	0	0	0
Lily	0	0	0	0

The resulting matrix shows us which vertices are exactly **two edges apart**.

An Interesting Property of Adjacency Matrices

Consider the following graph: And its associated A.M.:



	Joe	Mary	Tsuen	Lily
Joe	0	1	0	0
Mary	0	0	1	0
Tsuen	0	0	0	1
Lily	0	1	0	0

	Joe	Mary	Tsuen	Lily
Joe	0	0	1	0
Mary	0	0	0	1
Tsuen	0	1	0	0
Lily	0	0	1	0

	Joe	Mary	Tsuen	Lily
Joe	0	0	0	1
Mary	0	0	0	0
Tsuen	0	0	0	0
Lily	0	0	0	0

	Joe	Mary	Tsuen	Lily
Joe	0	0	0	0
Mary	0	0	0	0
Tsuen	0	0	0	0
Lily	0	0	0	0

	Joe	Mary	Tsuen	Lily
Joe	0	0	0	0
Mary	0	0	0	0
Tsuen	0	0	0	0
Lily	0	0	0	0

And if we multiply our new matrix by the original matrix again, we'll get all vertices that are exactly **3 edges apart**!

Another Way to Represent a Graph

Question:

How else can we represent a graph (without a 2D array)?

Answer:

A **directed graph** of n vertices can be represented by an array of n linked lists. This is called an **adjacency list**.

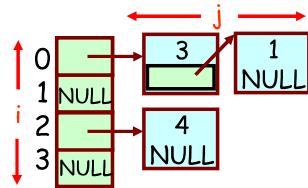
```
list<int> graph[n];
```

If we add a number j , to list number i (e.g., to $\text{graph}[i]$), this means that there is an edge from vertex i to vertex j .

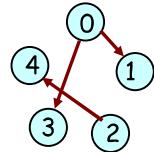
The Adjacency List

If we add a number j , to list number i (e.g., to $\text{graph}[i]$), this means that there is an edge from vertex i to vertex j

```
list<int> graph[4];
// edge from node 0 to node 3
graph[0].push_back(3);
graph[2].push_back(4);
graph[0].push_back(1);
```



So for each entry j , in list i , this means that there is an edge from vertex i to vertex j .



Which Representation Should You Use?

When should you use an **adjacency matrix** vs. an **adjacency list**?

Scenario #1:

We've got **10,000,000 users** who have relationships with each other - typically each person is friends with just a **few hundred** other people.

What would you do?

Option A: Store the graph in a **10 million by 10 million array**?
(That's **100 trillion** cells)



Option B: Store your graph in an array holding **10 million linked lists**, each holding roughly **500 items**?
(That's only **5 billion** pieces of data)

Which Representation Should You Use?

When should you use an **adjacency matrix** vs. an **adjacency list**?

Scenario #2:

We've got **1,000 cities**, with airlines offering flights from every city to almost every other city.

What would you do?

Option A: Store the graph in a **1000 by 1000 array**?
(That's **1 million** cells)

Option B: Store your graph in an array holding **1000 linked lists**, each holding roughly **1000 items**?
(That's also **1 million** pieces of data, but it's more complex)

Which Representation Should You Use?

When should you use an **adjacency matrix** vs. an **adjacency list**?

Use an **adjacency matrix** if you have **lots of edges** between vertices but **few vertices** ($< 10,000$ vertices).

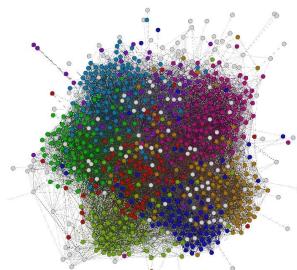
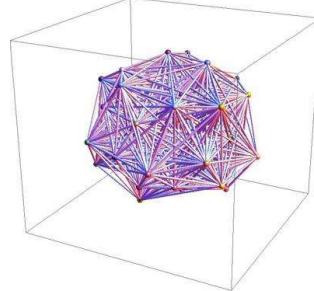
Use an **adjacency list** if you have **few edges** between vertices and lots of vertices ($> 10,000$ vertices).

A graph that has **many edges** between the vertices is called a "**dense graph**".

A graph that has **few edges** between the vertices is called a "**sparse graph**".

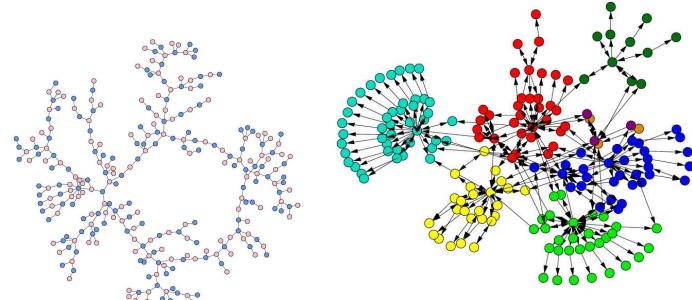
Let's see examples of both...

Dense(r) Graphs



Friendships on Facebook for people from Caltech.

Sparse Graphs



(High-school dating habits)

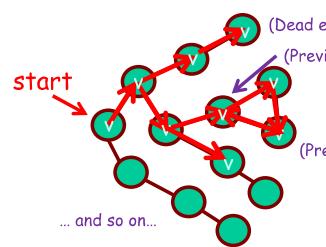
(Intra-website links)

Graph Traversals

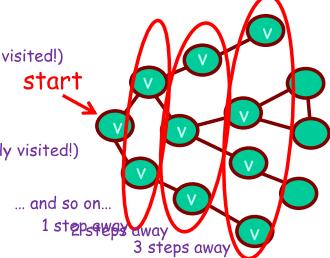
We can traverse graphs just like we traverse binary trees!

There are two types of graph traversals:
Depth-first and Breadth-first

A Depth-first Traversal keeps moving forward until it hits a **dead end** or a **previously-visited vertex**... then it backtrack and tries another path



A Breadth-first Traversal explores the graph in **growing concentric circles**, exploring all vertices 1 away from the start, then 2 away, then 3 away, etc.



Depth-first Traversals

Let's learn the **Depth-first Traversal** algorithm first:

Depth-First-Traversal(*curVertex*)

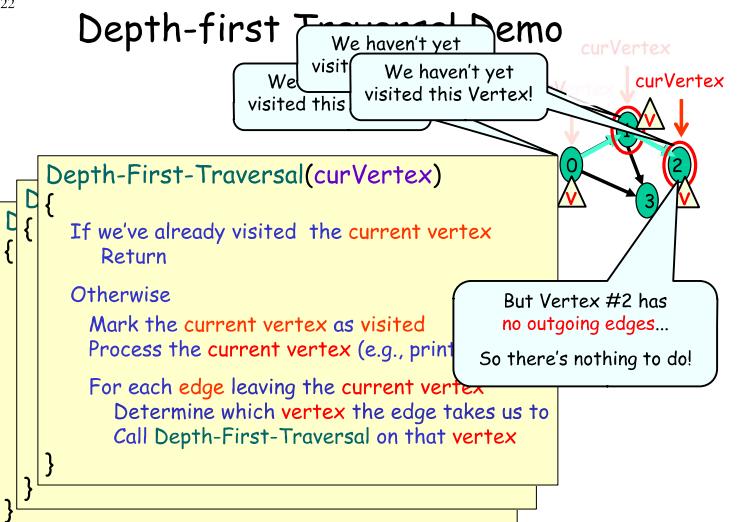
```
{
  If we've already visited the current vertex
    Return

  Otherwise
    Mark the current vertex as visited
    Process the current vertex (e.g., print it out)

    For each edge leaving the current vertex
      Determine which vertex the edge takes us to
      Call Depth-First-Traversal on that vertex
}
```

(Notice that it's recursive!)

Depth-first Traversal Demo



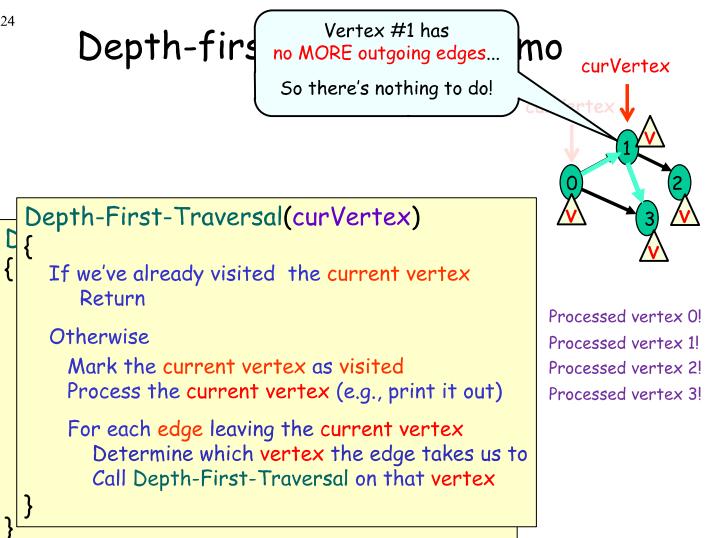
Depth-First-Traversal(*curVertex*)

```
{
  If we've already visited the current vertex
    Return

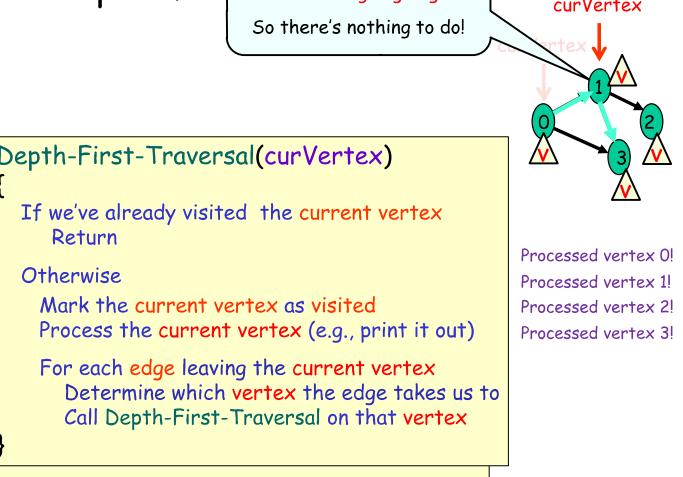
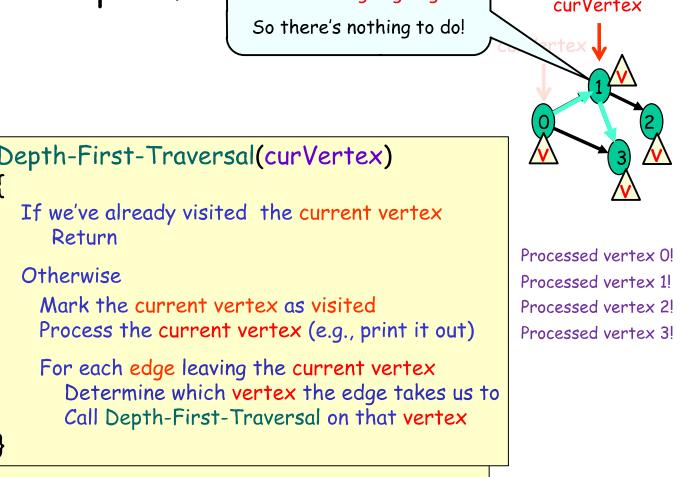
  Otherwise
    Mark the current vertex as visited
    Process the current vertex (e.g., print it out)

    For each edge leaving the current vertex
      Determine which vertex the edge takes us to
      Call Depth-First-Traversal on that vertex
}
```

Depth-first Traversal Demo



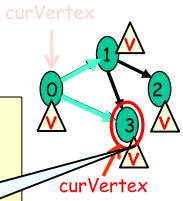
Depth-first Traversal Demo



Depth-first Traversal Demo

```

D { Depth-First-Traversal(curVertex)
  {
    If we've already visited the current vertex
      Return
    Otherwise
      Mark the current vertex
      Process the current vertex
      For each edge leaving the current vertex
        Determine which vertex the edge takes us to
        Call Depth-First-Traversal on that vertex
  }
}
  
```

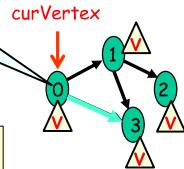


But we've already visited vertex #3!
So we don't want to do so again!

Processed vertex 0!
Processed vertex 1!
Processed vertex 2!
Processed vertex 3!

Depth-first Traversal Demo

Vertex #0 has no MORE outgoing edges...
So there's nothing to do!



```
Depth-First-Traversal(curVertex)
```

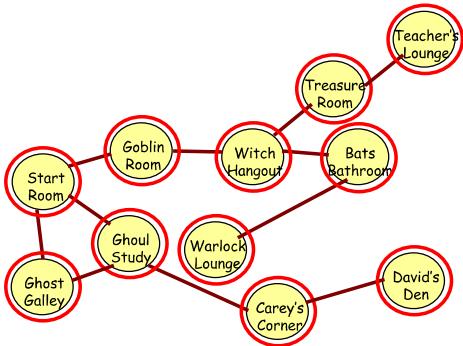
```
{
  If we've already visited the current vertex
  Return
  Otherwise
    Mark the current vertex as visited
    Process the current vertex (e.g., print it out)
    For each edge leaving the current vertex
      Determine which vertex the edge takes us to
      Call Depth-First-Traversal on that vertex
}
```

Processed vertex 0!
Processed vertex 1!
Processed vertex 2!
Processed vertex 3!

And we're done!

Depth-first Traversal Challenge

What does a Depth-first Traversal look like on this graph?



Implementing Depth-first Traversal w/ Stack!

You can also implement your Depth-first Traversal with a stack if you like! (What's not to like???)

```
Depth-First-Search-With-Stack(start_room)
```

```

Push start_room on the stack
While the stack is not empty
  Pop the top item off the stack and put it in variable c
  If c hasn't been visited yet
    Drop a breadcrumb (we've visited the current room)
    For each door d leaving the room
      If the room r behind door d hasn't been visited
        Push r onto the stack.
  
```

Basically, the stack allows you to simulate recursion...

Or does the recursion allow you to simulate a stack?

Hmmmmmmmm!

Breadth-first Graph Traversal

Idea:

Process all of the vertices that are 1 edge away from the start vertex,

then process all vertices that are two edges away,

then process all vertices that are three edges away,

etc...

Question:

What data structure could we use to implement this?

Answer:

Not a P, but a ?

Breadth-first Graph Traversal

```
Breadth-First-Search (startVertex)
```

```
{
  Add the starting vertex to our queue
  Mark the starting vertex as "discovered"
  While the queue is not empty
    Dequeue the top vertex from the queue and place in c
    Process vertex c (e.g., print its contents out)
    For each vertex v directly reachable from c
      If v has not yet been "discovered"
        Mark v as "discovered"
        Insert vertex v into the queue
}
```

Hmmm. Does this algorithm look familiar?

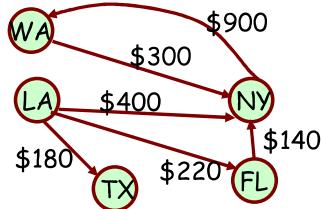
It's a-maze-ingly similar to our queue-based maze-solving algorithm!!!

Graphs With Weighted Edges

What does it mean for a graph to have **weighted edges**?

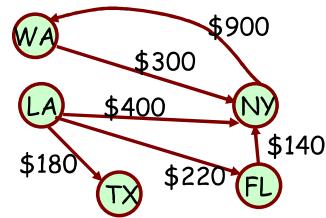
Definition: Each **edge** connecting **vertex u** with **vertex v** has a **weight** or **cost** associated with it.

Question: Why would we want to have weighted edges?



Graphs With Weighted Edges

Definition: The **weight** of a path from **vertex u** to **vertex v** is the **sum of the weights of the edges** between the two vertices.



Question: What's the cost of traveling from LA to NY to WA?

Question: What's the **shortest path** from LA to WA?

Definition: The **shortest path** between two vertices is the path with the lowest total cost of edges between the two vertices. (**The shortest path is a set of vertices**)

Finding the Shortest Path

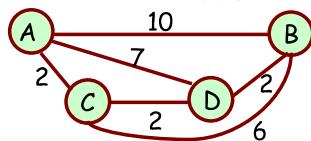
Question: How can we find the shortest path between any two nodes in a graph?

Answer: Dijkstra's Algorithm (the dorm guy?)

Dijkstra's Algorithm:

the length of

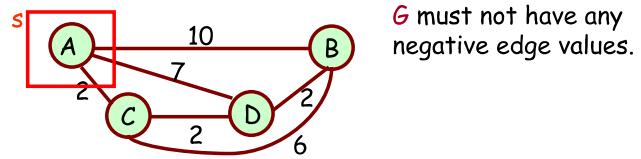
This algorithm determines the **shortest path** (i.e. set of vertices) from a start vertex **s** to all other vertices in the graph.



So Dijkstra(A) would give us a value of 6 for A to B, a value of 2 for A to C, and 4 for A to D.

Dijkstra's Algorithm

Input: A graph **G**, and a starting vertex **s**



G must not have any negative edge values.

Output: An array called **Dist** of optimal distances from **s** to every other node in the graph.

Dist	A	B	C	D
	0	6	2	4

Dijkstra's Algorithm: Basic Idea

Dijkstra's algorithm splits vertices in two distinct sets: the set of **unsettled** vertices and the set of **settled** vertices.

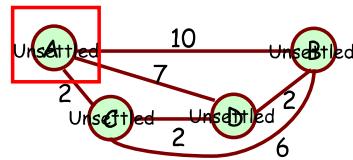
Unsettled vertex: A vertex **v** is **unsettled** if we don't know the optimal distance to it from the starting vertex **s**.

Settled vertex: A vertex **v** is **settled** if we have learned the optimal distance to it from the starting vertex **s**.

Initially all vertices are **unsettled**.

The algorithm ends once all vertices are in the **settled** set.

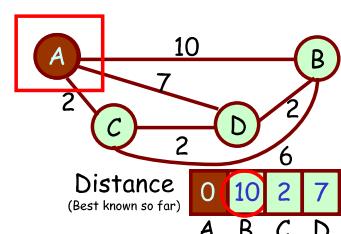
Start vertex



Dijkstra on a Graph

Assume that **all** vertices are **infinitely far away** to start...

Since we start at vertex **A**, we know it's the closest vertex to us! How far is it?
Zero steps away! We can settle it immediately!



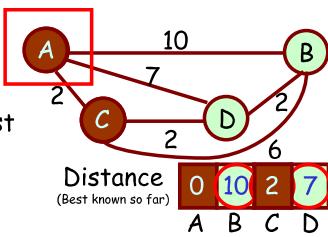
Now let's see which unsettled vertices we can reach directly from **A**.

- **B** is 10 units away.
- **C** is 2 units away.
- **D** is 7 units away.

And going directly from **A** to **D** is only 7 units away, which is less than infinity, so I'll update this entry too...

Dijkstra on a Graph

Ok, so now we know the costs to travel to all vertices directly reachable from **A**.



Which unsettled vertex is closest to **A**?

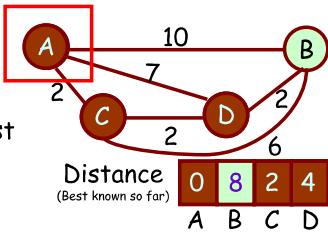
Right! **C** is closest to **A**.

If we go directly to **C** (**A** → **C**), it costs us 2 units. Is there any possible way I can travel to **C** cheaper by going through **B** or through **D** first?

So I know that if I travel directly from **A** to **C**, at a cost of 2 units, that's the *fastest* possible route. Therefore we can settle **C** at 2 units.

Dijkstra on a Graph

Ok, so now we know the best cost to get to all unsettled vertices, assuming we travel through **C**.



Which unsettled vertex is closest to **A** now?

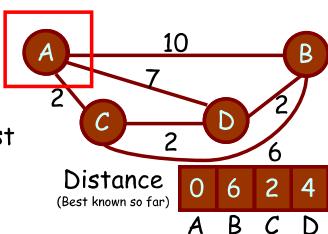
Right! **D** is closest.

If we take the path **A** → **C** → **D**, it costs us 4 units. Is there any possible way I can travel to **D** cheaper by going through **B** first?

So I know that if I travel from **A** → **C** → **D**, at a cost of 4 units, that's the *fastest* possible route. Therefore we can settle **D** at 4 units.

Dijkstra on a Graph

Ok, so now we know the best cost to get to all unsettled vertices, assuming we travel through **D**.



Which unsettled vertex is closest to **A** now?

Right! **B** is closest.

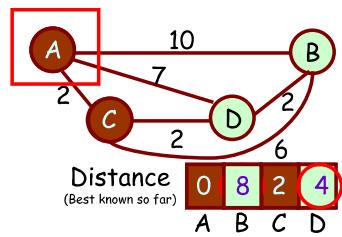
And now that all of our vertices are settled, we are guaranteed to have found the *minimum* travel distances to each of our vertices!

Dijkstra on a Graph

At this point, we know the shortest path from **A** to **C**. Now let's see if we can travel through **C** to reach one of our other unsettled vertices faster.

Ok, which unsettled vertices can be reached directly from **C**?

- **B** is 6 units away.
- **D** is 2 units away.



Let's do **D** next. We know we can get from **A** to **C** in 2 units, and we can directly go from **C** to **D** in 2 units, so we can reach **D** in just 4 units!

Is our new distance to **D** better than our old one?

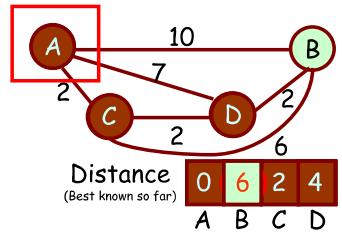
Yup!! Let's update our table again!

Dijkstra on a Graph

At this point, we know the shortest path from **A** to **D**. Now let's see if we can travel through **D** to reach one of our other unsettled vertices faster.

Ok, which unsettled vertices can be reached directly from **D**?

- **B** is 2 units away.



Let's check **B**. We know we can get from **A** to **D** in 4 units, and we can directly go from **D** to **B** in 2 units, so we can reach **B** in just 6 units!

Is our new distance to **B** better than our old one?

You bet!! Let's update our table!

Dijkstra

And now I'll give you the more formal algorithm...



Born: 11 May 1930, Rotterdam, Netherlands
Died: 6 August 2002, Nuenen, Netherlands

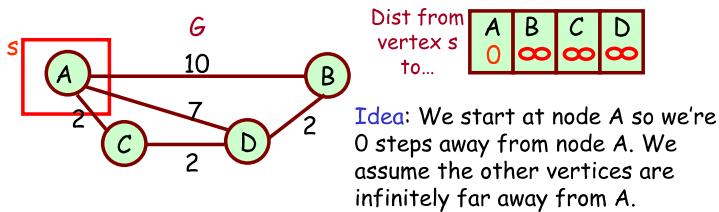
Dijkstra's Algorithm

Dijkstra's Algorithm uses 2 data structures:

- An array called **Dist** that holds the **the current best known cost** to get from **s** to every other vertex in the graph.

For each vertex **i**, **Dist[i]** starts out with a value of:

- 0** for vertex **s**
- Infinity** for all other vertices

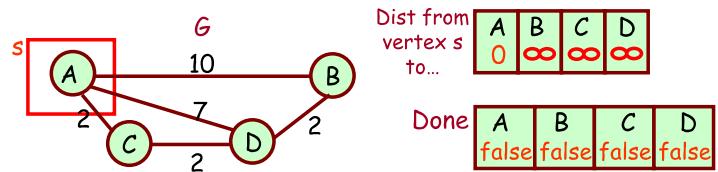


Dijkstra's Algorithm

Dijkstra's Algorithm uses 2 data structures:

- An array called **Done** that holds **true** for each vertex that has been fully processed, and **false** otherwise.

For each vertex **i**, **Done[i]** starts out with a value of **false**.



Dijkstra's Algorithm

While there are still unprocessed vertices:

Set **u** = the closest unprocessed vertex to the start vertex **s**

Mark vertex **u** as processed: **Done[u] = true**.

We now know how to reach **u** optimally from **s**

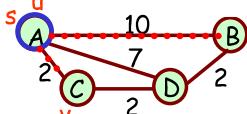
Loop through all unprocessed vertices:

Set **v** = the next unprocessed vertex

If there's an edge from **u** to **v** then compare:

- the previously computed path from **s** to **v** (i.e. **Dist[v]**) OR
- the path from **s** to **u**, and then from **u** to **v** (I.e. **Dist[u] + weight(u,v)**)

If the new cost is less than old cost then Set **Dist[v] = Dist[u] + weight(u,v)**



Dist from vertex s to...

A	B	C	D
0	10	2	∞

Done

U	V
true	false

Dijkstra's Algorithm

While there are still unprocessed vertices:

Set **u** = the closest unprocessed vertex to the start vertex **s**

Mark vertex **u** as processed: **Done[u] = true**.

We now know how to reach **u** optimally from **s**

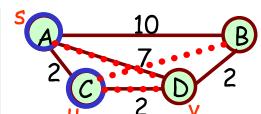
Loop through all unprocessed vertices:

Set **v** = the next unprocessed vertex

If there's an edge from **u** to **v** then compare:

- the previously computed path from **s** to **v** (i.e. **Dist[v]**) OR
- the path from **s** to **u**, and then from **u** to **v** (I.e. **Dist[u] + weight(u,v)**)

If the new cost is less than old cost then Set **Dist[v] = Dist[u] + weight(u,v)**



Dist from vertex s to...

A	B	C	D
0	10	2	4

Done

U	V
true	false

Dijkstra's Algorithm

While there are still unprocessed vertices:

Set **u** = the closest unprocessed vertex to the start vertex **s**

Mark vertex **u** as processed: **Done[u] = true**.

We now know how to reach **u** optimally from **s**

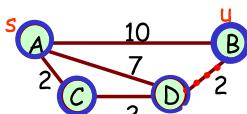
Loop through all unprocessed vertices:

Set **v** = the next unprocessed vertex

If there's an edge from **u** to **v** then compare:

- the previously computed path from **s** to **v** (i.e. **Dist[v]**) OR
- the path from **s** to **u**, and then from **u** to **v** (I.e. **Dist[u] + weight(u,v)**)

If the new cost is less than old cost then Set **Dist[v] = Dist[u] + weight(u,v)**



Dist from vertex s to...

A	B	C	D
0	6	2	4

Done

U
true

And we're done! The **Dist** array contains the results.